


# Binary Classification of Graphene Flakes



Samuel Rager and Alvin  
Wang



# Main References

- VII. Krizhevsky, Sutskever, & Hinton. "ImageNet classification with deep convolutional neural networks." *Commun. ACM* 60. 2017, accessed 29 October 2024, <[https://proceedings.neurips.cc/paper\\_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf)>.
- VIII. Kingma, Diederik P. "Adam: A method for stochastic optimization." *arXiv preprint arXiv:1412.6980* (2014), accessed 29 October 2024, <<https://arxiv.org/pdf/1412.6980>>.
- XI. He, Kaiming, et al. "Deep residual learning for image recognition." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016, accessed 29 October 2024, <<https://arxiv.org/pdf/1512.03385v1>>.


# Data Cleaning Methodology



# Timeline

2010	2012	2014	2015	2016	2017	2018	2019	2023
•First ImageNet challenge	•AlexNet	•Adam •VGG •GoogleNet	•ResNet •Batch Norm. •GoogleNet renamed Inception •Yolo V1	•Inception v3	•AdamW •Transformer •Last Imagenet Challenge	•Mobile Net v2 (Inverted Residual)	•Efficient Net	•Yolo V8

# AlexNet






# Convolution



# Max Pooling

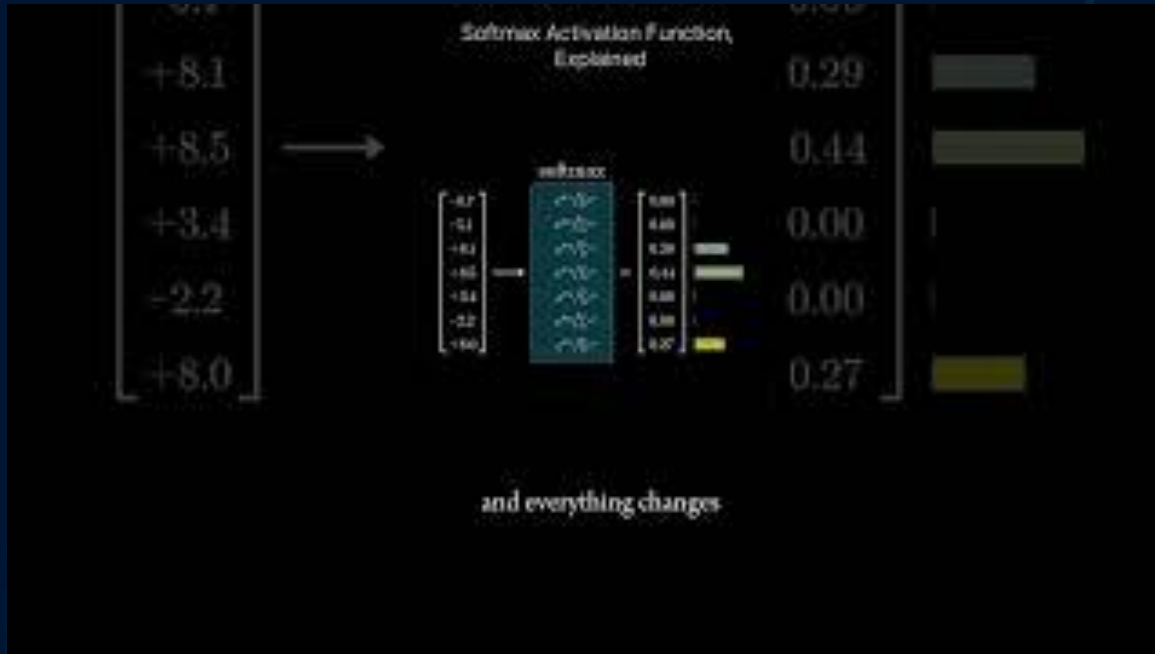


# Rectified Linear Function






# Softmax- Reduces to Sigmoid in Binary Case



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



# Local Response Normalization

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

**For AlexNet, k=2, n=5, alpha=10<sup>-4</sup>,  
beta=0.75**

# Data Structure for Colab Models

Data Structure

Dataset-|

| \_>train---| \_> graphene  
|\_> non\_graphene

| \_>test---| \_> graphene  
|\_> non\_graphene

# Adam (Adaptive Moment Estimation/Gradient Descent with Momentum)

**Algorithm 1:** *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation.  $g_t^2$  indicates the elementwise square  $g_t \odot g_t$ . Good default settings for the tested machine learning problems are  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-8}$ . All operations on vectors are element-wise. With  $\beta_1^t$  and  $\beta_2^t$  we denote  $\beta_1$  and  $\beta_2$  to the power  $t$ .

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1)$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

**end while**

**return**  $\theta_t$  (Resulting parameters)



# Backpropagation

## Algorithm 1 Backpropagation

```
1: procedure BACKPROPAGATION( $X, \theta$ )
2:    $a^0 \leftarrow x$  ▷ Set Input
3:   for  $l \leftarrow [1, L]$  do ▷ Forward Pass: Compute weighted inputs and activations
4:      $z^l \leftarrow W^l a^{l-1} + b^l$ 
5:      $a^l \leftarrow \sigma(z^l)$ 
6:   end for
7:
8:    $\delta_L \leftarrow (a^L - y) \odot \sigma'(z^L)$  ▷ Compute Error in Final Layer
9:   for  $l \leftarrow [L - 1, 1]$  do ▷ Backpropagate Error
10:     $\delta^l \leftarrow (W^{l+1})^T \delta^{l+1} \odot \sigma'(z^l)$ 
11:  end for
12:
13:  for  $l \leftarrow [1, L]$  do ▷ Compute gradients w.r.t. network parameters
14:     $\nabla_{b^l} = \delta^l$ 
15:     $\nabla_{W^l} = \delta^l (a^{l-1})^T$ 
16:  end for
17:  return  $\nabla_{\theta} = (\nabla_{W^l}, \nabla_{b^l})$ 
18: end procedure
```

# Binary Cross Entropy Loss

$$L = -\frac{1}{N} \sum_{i=1}^N [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)]$$

➤  $N$  is the number of samples

➤  $y_i$  is the true label for the  $i^{\text{th}}$  sample (0 or 1)

➤  $p_i$  is the predicted probability that the  $i^{\text{th}}$  sample belongs to Positive class




## Binary Cross Entropy with Logits Loss


$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N \left[ \max(z_i, 0) - z_i y_i + \log \left( 1 + e^{-|z_i|} \right) \right]$$

# AlexNet

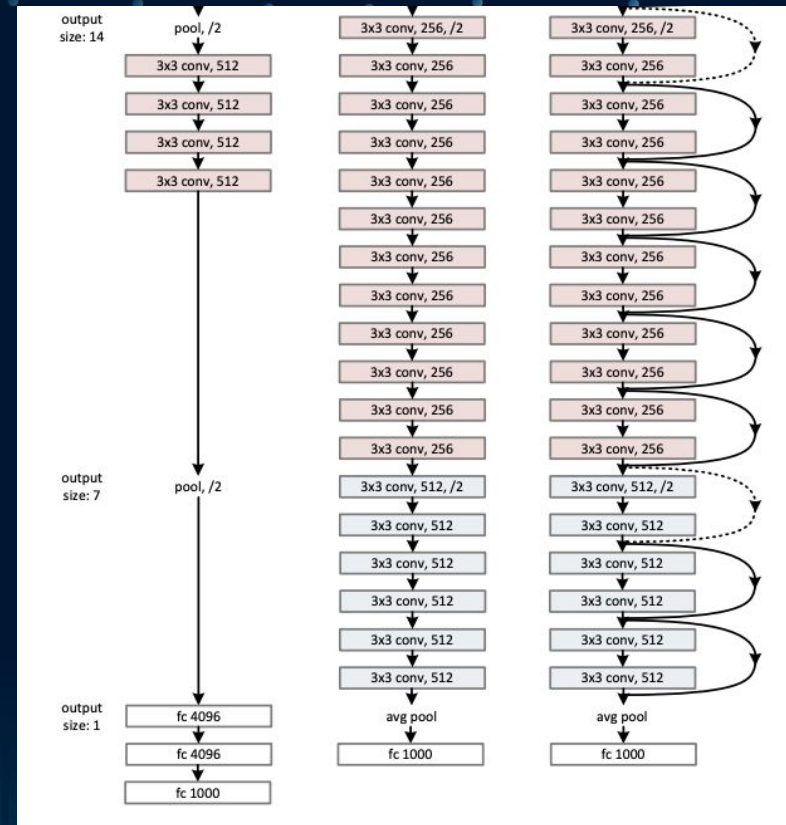
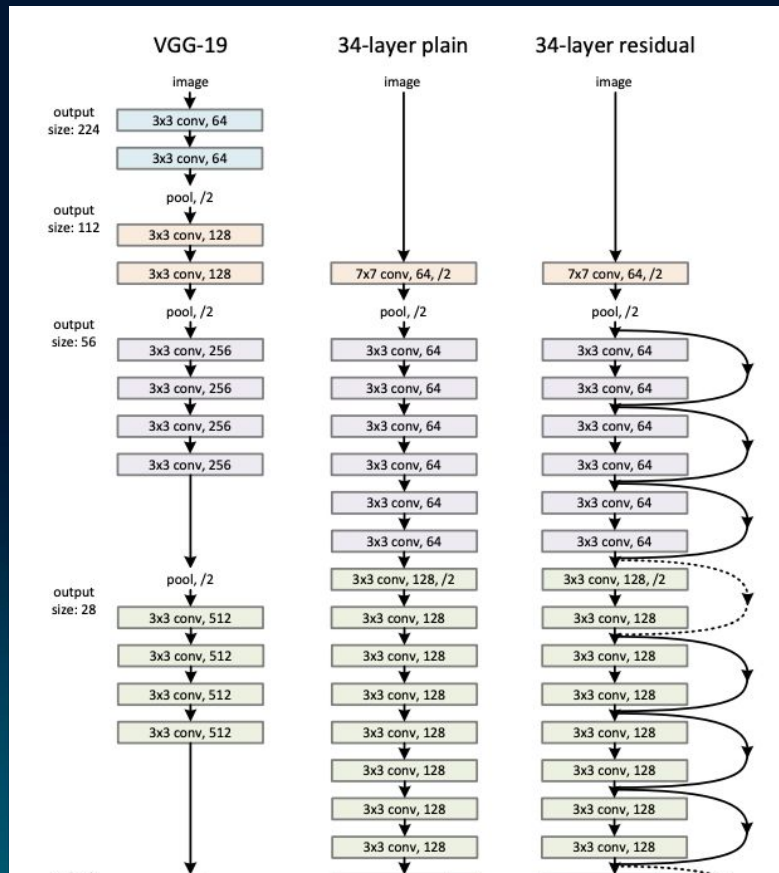
**Contribution Mnemonic:  
Programming working  
backwards from desired  
output**



# Visual Geometry Group (VGG) 16



# ResNet Architecture





# ResNet 50 Architecture

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	$112 \times 112$	$7 \times 7$ , 64, stride 2				
conv2_x	$56 \times 56$	$3 \times 3$ max pool, stride 2				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	$28 \times 28$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	$14 \times 14$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	$7 \times 7$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	$1 \times 1$	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

# Batch Normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

**Input:** Network  $N$  with trainable parameters  $\Theta$ ;  
subset of activations  $\{x^{(k)}\}_{k=1}^K$


**Output:** Batch-normalized network for inference,  $N_{\text{BN}}^{\text{inf}}$

- 1:  $N_{\text{BN}}^{\text{tr}} \leftarrow N$  // Training BN network
- 2: **for**  $k = 1 \dots K$  **do**
- 3:   Add transformation  $y^{(k)} = \text{BN}_{\gamma^{(k)}, \beta^{(k)}}(x^{(k)})$  to  $N_{\text{BN}}^{\text{tr}}$  (Alg. 1)
- 4:   Modify each layer in  $N_{\text{BN}}^{\text{tr}}$  with input  $x^{(k)}$  to take  $y^{(k)}$  instead
- 5: **end for**
- 6: Train  $N_{\text{BN}}^{\text{tr}}$  to optimize the parameters  $\Theta \cup \{\gamma^{(k)}, \beta^{(k)}\}_{k=1}^K$
- 7:  $N_{\text{BN}}^{\text{inf}} \leftarrow N_{\text{BN}}^{\text{tr}}$  // Inference BN network with frozen parameters
- 8: **for**  $k = 1 \dots K$  **do**
- 9:   // For clarity,  $x \equiv x^{(k)}, \gamma \equiv \gamma^{(k)}, \mu_{\mathcal{B}} \equiv \mu_{\mathcal{B}}^{(k)}$ , etc.
- 10:   Process multiple training mini-batches  $\mathcal{B}$ , each of size  $m$ , and average over them:  
$$\text{E}[x] \leftarrow \text{E}_{\mathcal{B}}[\mu_{\mathcal{B}}]$$
$$\text{Var}[x] \leftarrow \frac{m}{m-1} \text{E}_{\mathcal{B}}[\sigma_{\mathcal{B}}^2]$$
- 11:   In  $N_{\text{BN}}^{\text{inf}}$ , replace the transform  $y = \text{BN}_{\gamma, \beta}(x)$  with  
$$y = \frac{\gamma}{\sqrt{\text{Var}[x] + \epsilon}} \cdot x + \left(\beta - \frac{\gamma \text{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}}\right)$$
- 12: **end for**

**Algorithm 2:** Training a Batch-Normalized Network

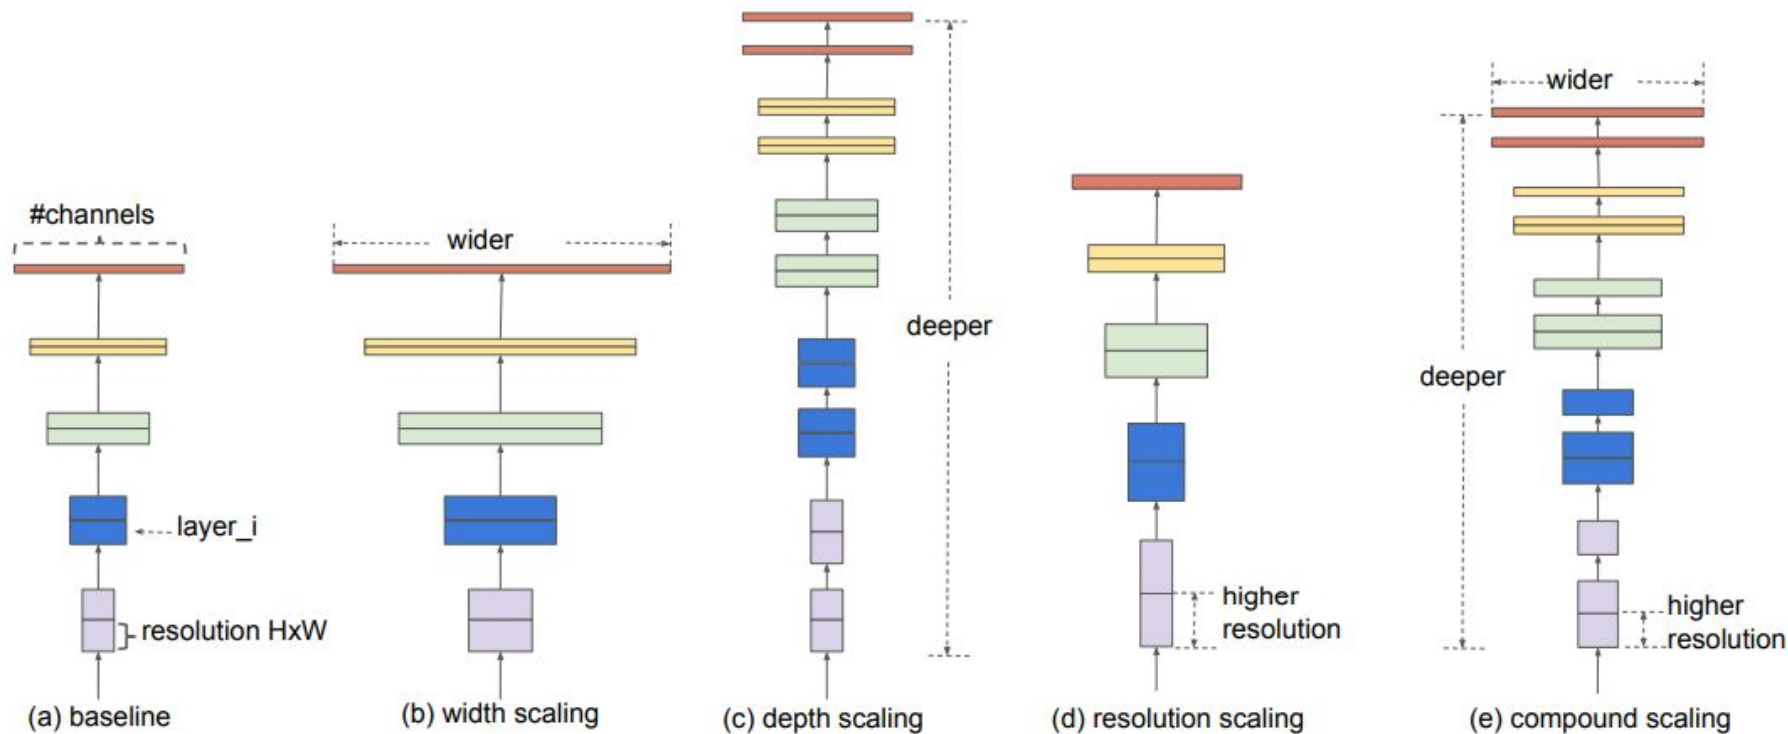


# EfficientNet B0



# EfficientNet B0

XIV.



# Grid Search




# EfficientNet B0 Architecture


Stage $i$	Operator $\hat{\mathcal{F}}_i$	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels $\hat{C}_i$	#Layers $\hat{L}_i$
1	Conv3x3	$224 \times 224$	32	1
2	MBConv1, k3x3	$112 \times 112$	16	1
3	MBConv6, k3x3	$112 \times 112$	24	2
4	MBConv6, k5x5	$56 \times 56$	40	2
5	MBConv6, k3x3	$28 \times 28$	80	3
6	MBConv6, k5x5	$14 \times 14$	112	3
7	MBConv6, k5x5	$14 \times 14$	192	4
8	MBConv6, k3x3	$7 \times 7$	320	1
9	Conv1x1 & Pooling & FC	$7 \times 7$	1280	1

# Inverted Residual Block


(a) Regular




(b) Separable



(c) Separable with linear bottleneck




(d) Bottleneck with expansion layer




# Inverted Residual Block

(a) Residual block

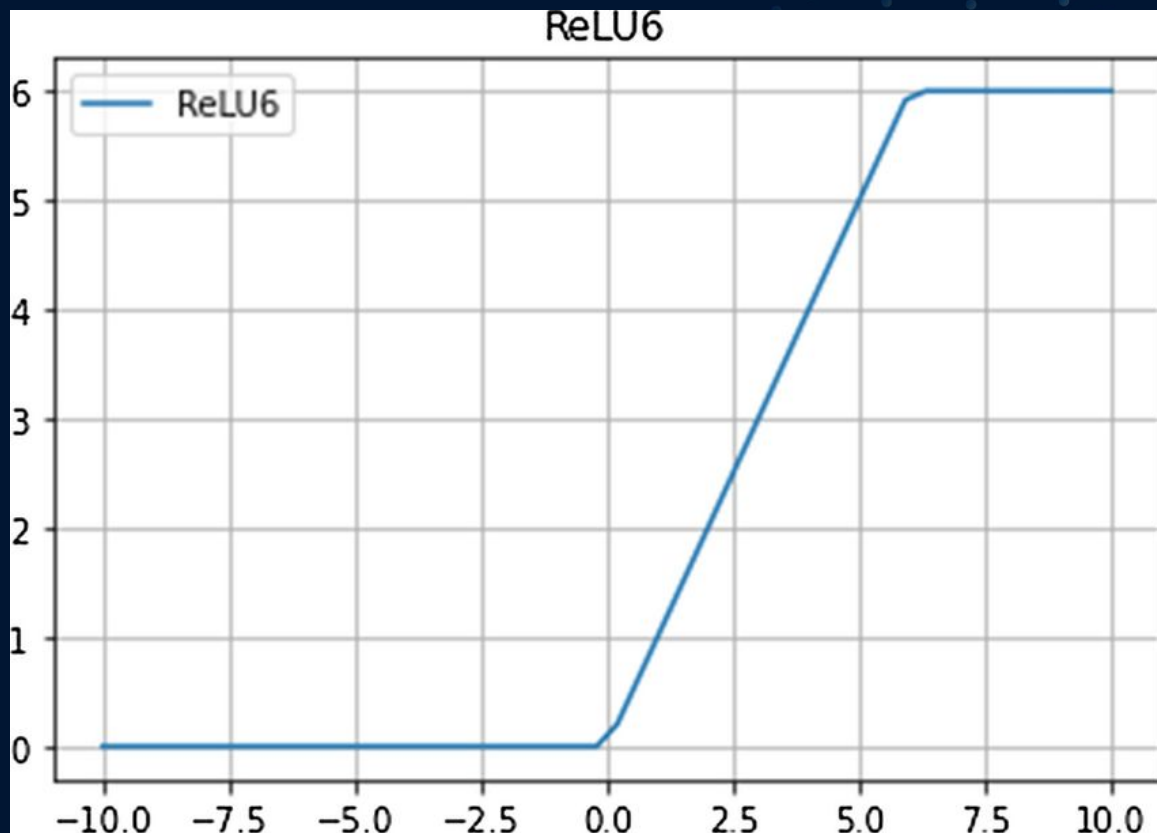


(b) Inverted residual block





# ReLU6



# Compound Scaling

Baseline ResNet-50 (He et al., 2016)	4.1B	76.0%
Scale ResNet-50 by depth ( $d=4$ )	16.2B	78.1%
Scale ResNet-50 by width ( $w=2$ )	14.7B	77.7%
Scale ResNet-50 by resolution ( $r=2$ )	16.4B	77.5%
<b>ResNet-50 compound scale</b>	<b>16.7B</b>	<b>78.8%</b>

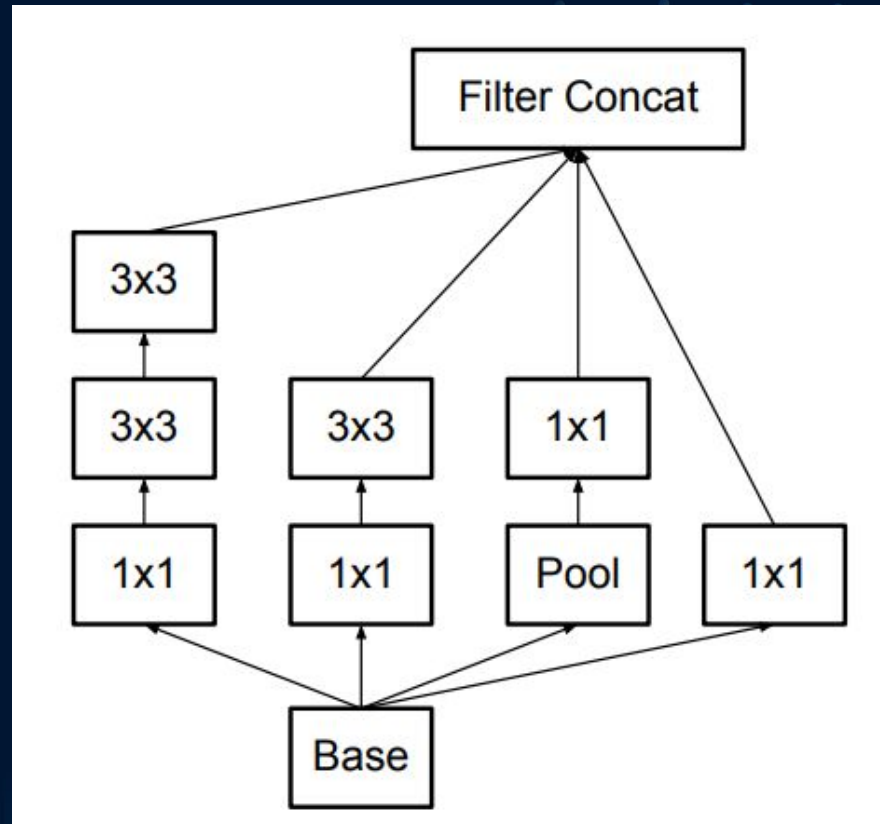
## Inception (GoogleNet)



# Inception v3 Architecture

type	patch size/stride or remarks	input size
conv	$3 \times 3 / 2$	$299 \times 299 \times 3$
conv	$3 \times 3 / 1$	$149 \times 149 \times 32$
conv padded	$3 \times 3 / 1$	$147 \times 147 \times 32$
pool	$3 \times 3 / 2$	$147 \times 147 \times 64$
conv	$3 \times 3 / 1$	$73 \times 73 \times 64$
conv	$3 \times 3 / 2$	$71 \times 71 \times 80$
conv	$3 \times 3 / 1$	$35 \times 35 \times 192$
$3 \times$ Inception	As in figure 5	$35 \times 35 \times 288$
$5 \times$ Inception	As in figure 6	$17 \times 17 \times 768$
$2 \times$ Inception	As in figure 7	$8 \times 8 \times 1280$
pool	$8 \times 8$	$8 \times 8 \times 2048$
linear	logits	$1 \times 1 \times 2048$
softmax	classifier	$1 \times 1 \times 1000$

**Figure 5**



**Figure 6**


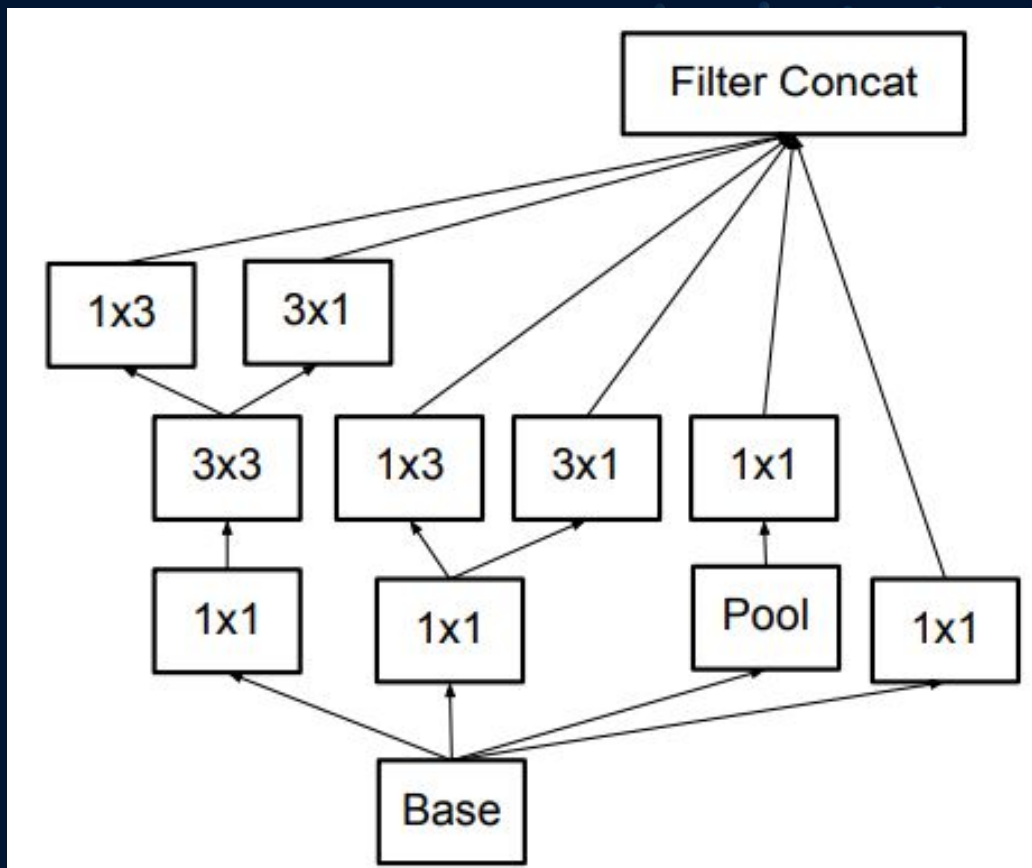






Figure 7



# Loss Per Epoch w/ Curve Fit for All Models





# Pre-Trained Models Used & Metrics

Metric AlexNet VGG16 ResNet50 EfficientNet b0 Inception v3						
0	Accuracy	0.98	0.98	0.98	0.99	0.99
1	Precision	0.98	0.98	1.00	0.99	0.99
2	F1	0.99	0.99	0.99	1.00	1.00
3	ROC AUC	0.50	0.40	0.96	1.00	0.97
4	Runtime (sec)	173.46	676.63	484.59	301.84	662.06

Metric	AlexNet	VGG16	ResNet50	EfficeintNet b0	Inception v3
Accuracy	1	0.93	0.89	0.83	0.93
Precision	1	0.89	0.83	0.75	0.88
F1	1	0.93	0.9	0.85	0.93
AUC of ROC	1	0.94	0.94	1	1
Runtime(sec)	118.23	244.61	596.17	358.95	724.81
Parameter #	61,100,840	138,357,544	25,557,032	5.3M	27,161,264



# AlexNet Confusion Matrix





# AlexNet ROC/ AUC Plot

$$TPR = \frac{TP}{TP + FN}$$

$$FPR = \frac{FP}{FP + TN}$$




# ResNet50 Confusion Matrix






# ResNet50 ROC/ AUC Plot



Receiver Operating Characteristic (ROC) Curve





Receiver Operating Characteristic (ROC) Curve





# VGG16 Confusion Matrix





# VGG16 ROC/ AUC Plot





# EfficientNet b0 Confusion Matrix



# EfficientNet b0 ROC/ AUC Plot





# Inception v2 Confusion Matrix





# Inception v2 ROC/ AUC Plot



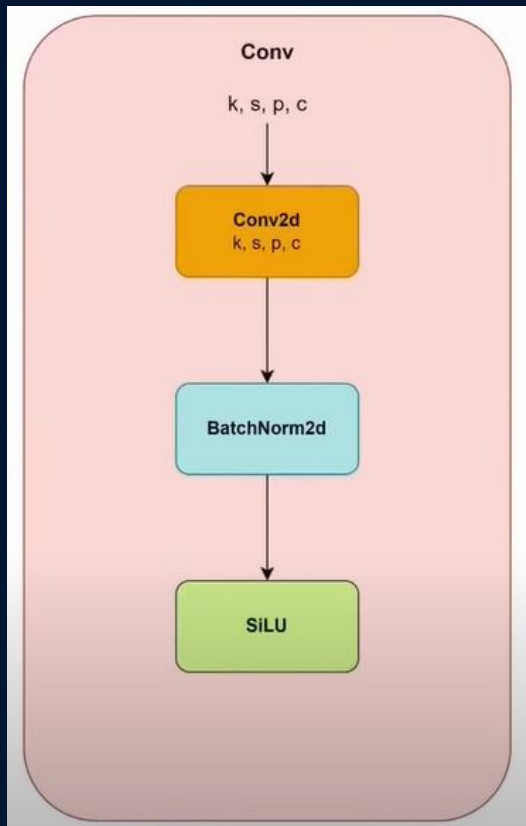
# YOLO Colab Implementation

```
!pip install ultralytics
from ultralytics import YOLO
model = YOLO("yolov8n.pt")
# Train the model
model.train(data='graphene.yaml', epochs=10, imgsz=640)
```

## Data Structure


```
|-->train
      |-->images
      |-->labels (.txt)
|-->val
      |-->images
      |-->labels (.txt)
```

# YOLOv8 Architecture




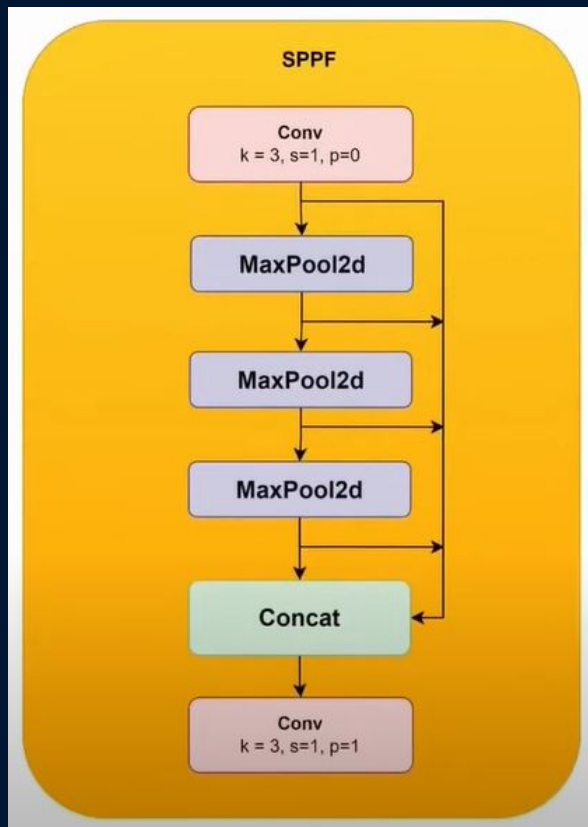
XXII.

$$\text{SiLU}(x) = x \left( \frac{1}{1 + e^{-x}} \right)$$

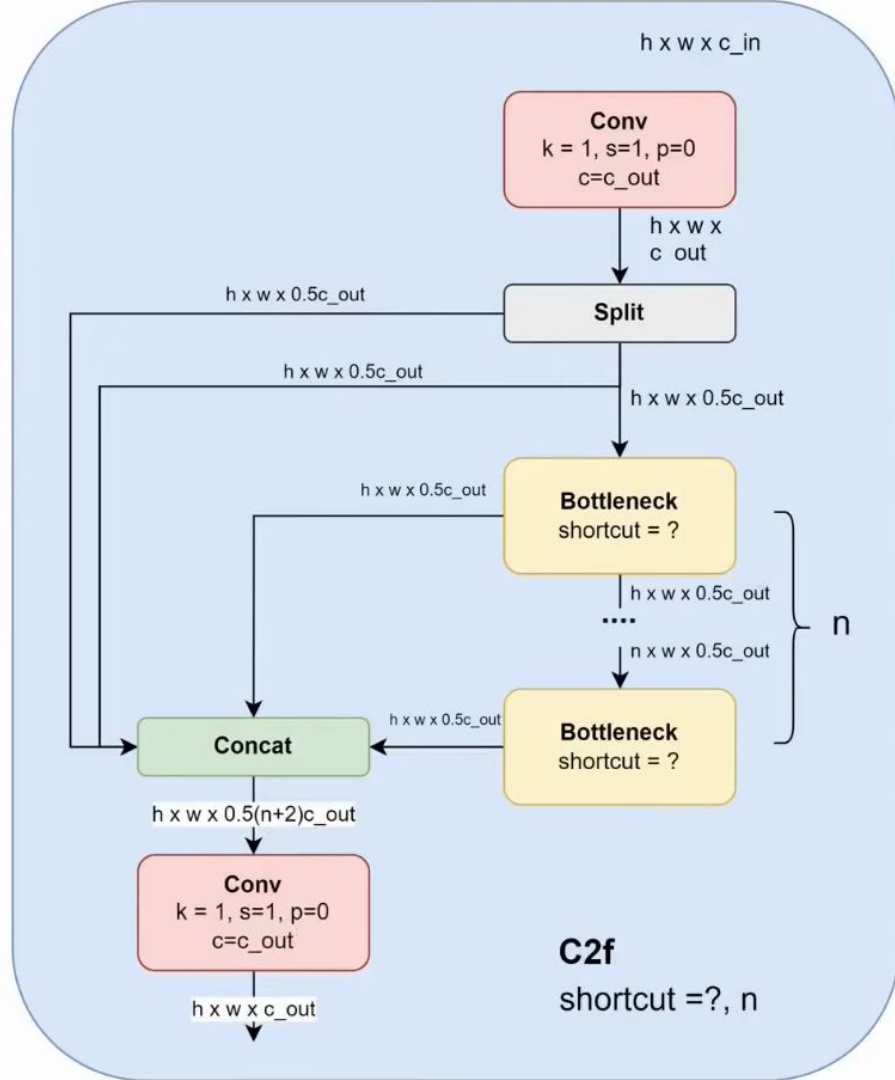


XXI.

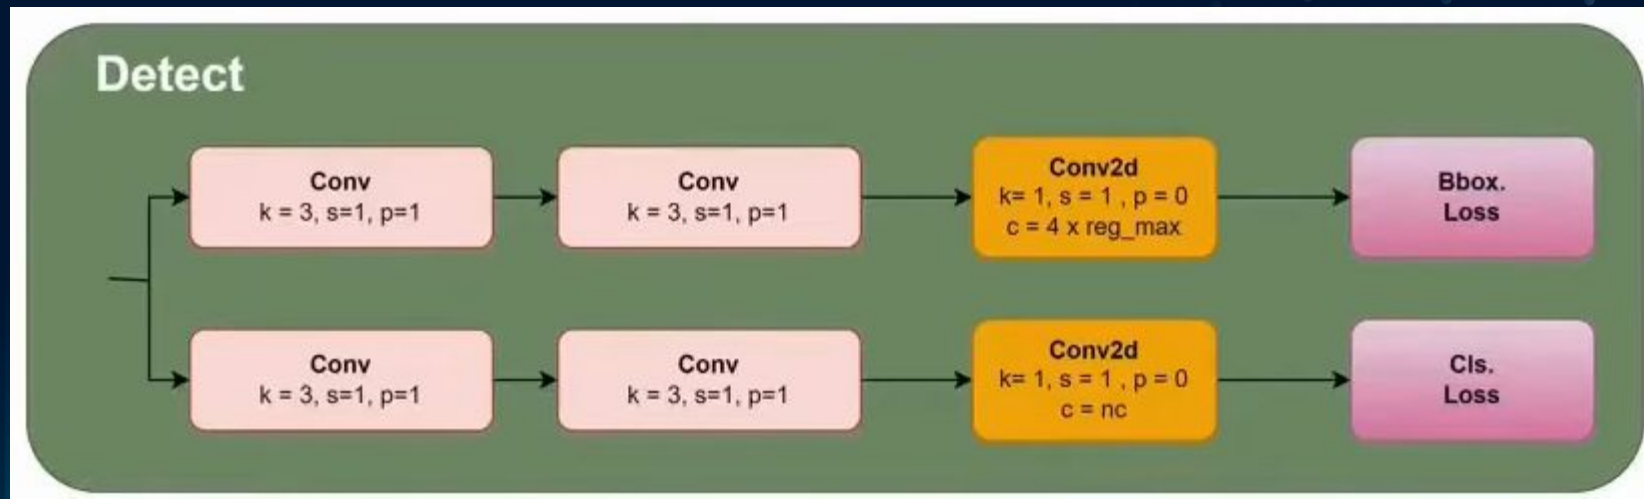
# YOLOv8 Architecture



# YOLOv8 Architecture



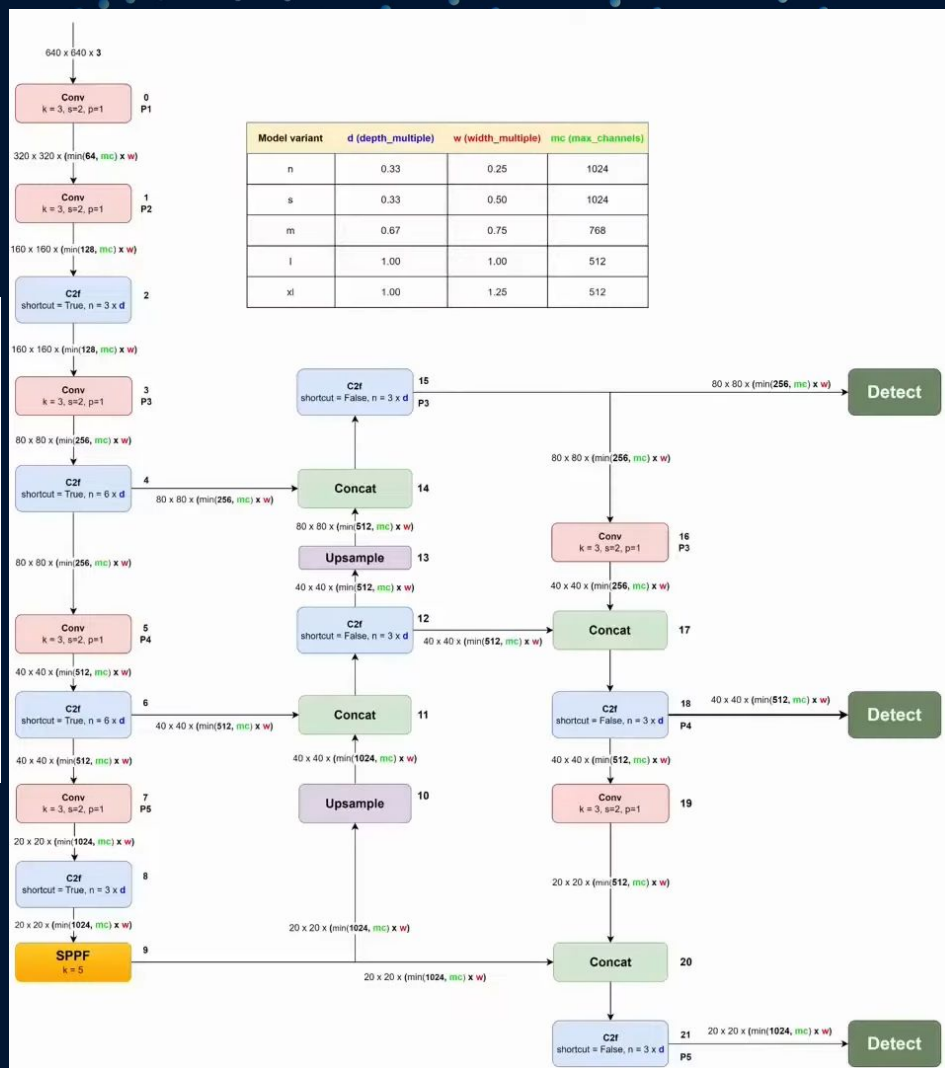
# YOLOv8 Architecture



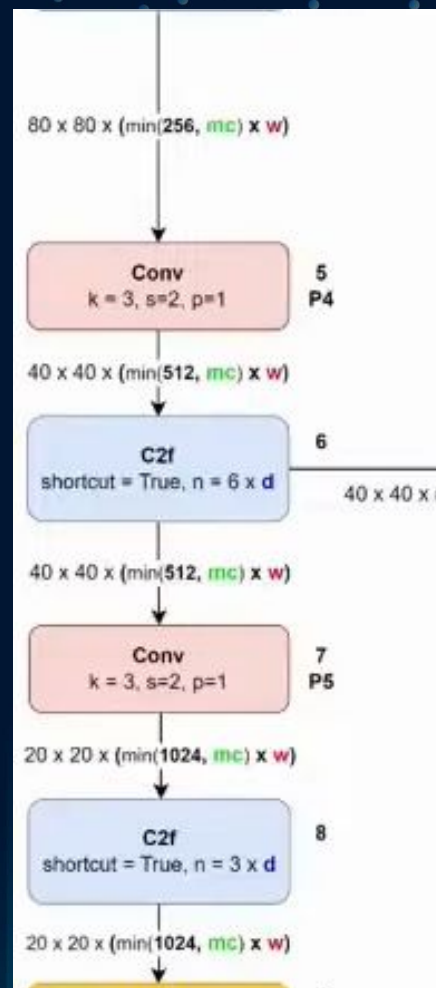
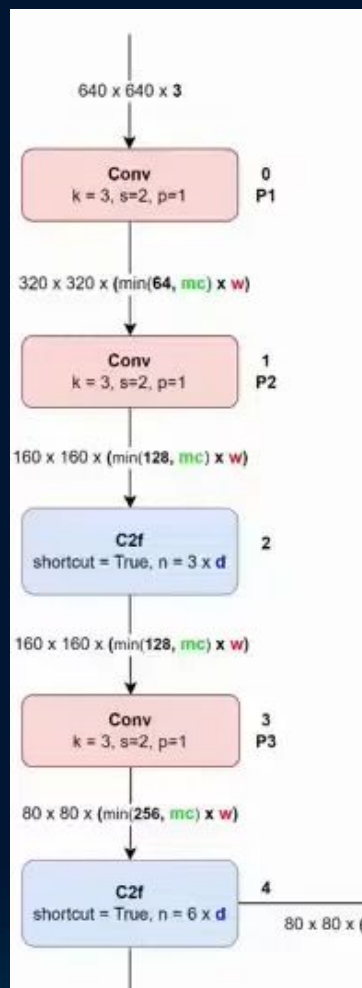


# YOLOv8 Architecture

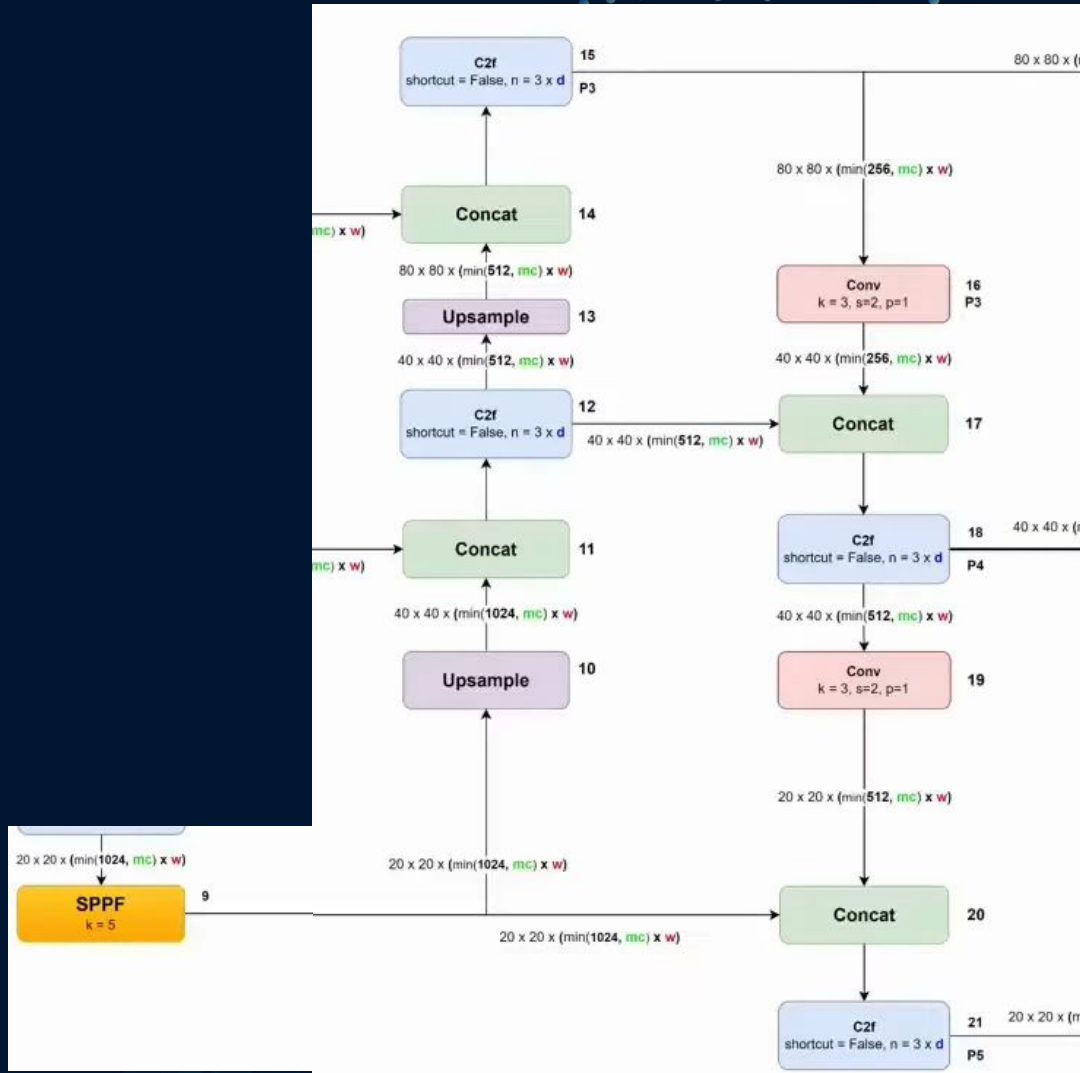
Model variant	d (depth_multiple)	w (width_multiple)	mc (max_channels)
n	0.33	0.25	1024
s	0.33	0.50	1024
m	0.67	0.75	768
l	1.00	1.00	512
xl	1.00	1.25	512



# Backbone



# Neck



# Head

80 x 80 x (min(256,  $mc$ ) x  $w$ )

Detect

40 x 40 x (min(512,  $mc$ ) x  $w$ )

Detect

20 x 20 x (min(1024,  $mc$ ) x  $w$ )

Detect

# AdamW Optimizer Psuedocode (Algo. 1 is

## Algorithm 2 Adam with $L_2$ regularization and Adam with decoupled weight decay (AdamW)

- 1: **given**  $\alpha = 0.001, \beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}, \lambda \in \mathbb{R}$
- 2: **initialize** time step  $t \leftarrow 0$ , parameter vector  $\theta_{t=0} \in \mathbb{R}^n$ , first moment vector  $m_{t=0} \leftarrow \mathbf{0}$ , second moment vector  $v_{t=0} \leftarrow \mathbf{0}$ , schedule multiplier  $\eta_{t=0} \in \mathbb{R}$
- 3: **repeat**
- 4:    $t \leftarrow t + 1$
- 5:    $\nabla f_t(\theta_{t-1}) \leftarrow \text{SelectBatch}(\theta_{t-1})$  ▷ select batch and return the corresponding gradient
- 6:    $g_t \leftarrow \nabla f_t(\theta_{t-1}) + \lambda \theta_{t-1}$
- 7:    $m_t \leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t$  ▷ here and below all operations are element-wise
- 8:    $v_t \leftarrow \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$
- 9:    $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  ▷  $\beta_1$  is taken to the power of  $t$
- 10:    $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  ▷  $\beta_2$  is taken to the power of  $t$
- 11:    $\eta_t \leftarrow \text{SetScheduleMultiplier}(t)$  ▷ can be fixed, decay, or also be used for warm restarts
- 12:    $\theta_t \leftarrow \theta_{t-1} - \eta_t \left( \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon) + \lambda \theta_{t-1} \right)$
- 13: **until** *stopping criterion is met*
- 14: **return** optimized parameters  $\theta_t$

# Acknowledgements

- Abhay Narayan Pasupathy Research Group <https://anp-lab.physics.columbia.edu/>
- Reproducibility: <https://github.com/56sarager/Graphene-Classification> ,  
[https://drive.google.com/drive/folders/1CJKs\\_0wGzySyGWV78xK-VB7FZPO\\_-qDO](https://drive.google.com/drive/folders/1CJKs_0wGzySyGWV78xK-VB7FZPO_-qDO)
- Johannes Geurs <https://www.nature.com/articles/s41598-018-23922-1> (paper develops a program for identifying graphene flake from microscope images like originally desired)
- "PYTORCH Documentation." *PyTorch Documentation - PyTorch 2.5 Documentation*, [pytorch.org/docs/stable/index.html](https://pytorch.org/docs/stable/index.html). Accessed 30 Oct. 2024.



# References

- I. Askari, Syed Sajjad. "Alex-Net Explanation and Implementation in Tensorflow and Keras." *Medium*, Medium, 14 Jan. 2023, [medium.com/@syedsajjad62/alex-net-explanation-and-implementation-in-tensorflow-and-keras-8047efeb7a0f](https://medium.com/@syedsajjad62/alex-net-explanation-and-implementation-in-tensorflow-and-keras-8047efeb7a0f).
- II. Sanderson, Grant. "But What is a Convolution?" *Youtube*, 3blue1brown, 29 Dec. 2023, <https://www.youtube.com/shorts/kpG7l2MOcnI>.
- III. Ave Coders. "101 Concepts of Data Science and Machine Learning" *Youtube*, AveCoders, 28 May. 2024, <https://www.youtube.com/shorts/ykt5PUxs3z8>.
- IV. Chima, Precious. "Activation Functions: ReLU & Softmax." *Medium*, Medium, 5 Apr. 2020, <https://medium.com/@preshchima/activation-functions-relu-softmax-87145bf39288>.
- V. Alexander, Giffah. "Softmax Activation function explained" *Youtube*, Giffah\_Alexander, 22 Oct. 2024, [https://www.youtube.com/shorts/SrJN\\_hpiuAs](https://www.youtube.com/shorts/SrJN_hpiuAs).
- VI. Sigmoid function. "Sigmoid function." *Wikipedia*, Wikipedia, 14 Nov. 2024, [https://en.wikipedia.org/wiki/Sigmoid\\_function](https://en.wikipedia.org/wiki/Sigmoid_function).
- VII. Krizhevsky, Sutskever, & Hinton. "ImageNet classification with deep convolutional neural networks." *Commun. ACM* 60. 2017, accessed 29 October 2024, <[https://proceedings.neurips.cc/paper\\_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf)>.

# References

- VIII. Kingma, Diederik P. "Adam: A method for stochastic optimization." *arXiv preprint arXiv:1412.6980* (2014), accessed 29 October 2024, <<https://arxiv.org/pdf/1412.6980>>.
- IX. Gabbard & Miller. "Machine Learning from Scratch: Stochastic Gradient Descent and Adam Optimizer." *18.0851 Project*, accessed 17 November 2024, <[https://www.mit.edu/~jgabbard/assets/180851\\_Project\\_final.pdf](https://www.mit.edu/~jgabbard/assets/180851_Project_final.pdf)>.
- X. Bangar, Siddhesh. "VGG-Net Architecture Explained." *Medium*, Medium, 28 Jun. 2022, <https://medium.com/@siddheshb008/vgg-net-architecture-explained-71179310050f>.
- XI. He, Kaiming, et al. "Deep residual learning for image recognition." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016, accessed 29 October 2024, <<https://arxiv.org/pdf/1512.03385v1>>.
- XII. Ioffe, Sergey. "Batch normalization: Accelerating deep network training by reducing internal covariate shift." *arXiv preprint arXiv:1502.03167* (2015). <<https://arxiv.org/pdf/1502.03167>>.
- XIII. Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." *arXiv preprint arXiv:1409.1556* (2014). <<https://arxiv.org/pdf/1409.1556>>.
- XIV. Tan, Mingxing, and Quoc Le. "Efficientnet: Rethinking model scaling for convolutional neural networks." *International conference on machine learning*. PMLR, 2019, accessed 29 October 2024, <<https://arxiv.org/pdf/1905.11946>>.

# References

- XV. Hien, Ngo Le Huy. "Grid Search space Illustration" *ResearchGate*, ResearchGate, [https://www.researchgate.net/figure/Grid-Search-space-illustration\\_fig9\\_346571789](https://www.researchgate.net/figure/Grid-Search-space-illustration_fig9_346571789).
- XVI. Sandler, Mark, et al. "Mobilenetv2: Inverted residuals and linear bottlenecks." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2018, <<https://arxiv.org/pdf/1801.04381>>.
- XVII. Neural Computing and Applications. "Graph for ReLU6 activation function" *ResearchGate*, ResearchGate, [https://www.researchgate.net/figure/Grid-Search-space-illustration\\_fig9\\_346571789](https://www.researchgate.net/figure/Grid-Search-space-illustration_fig9_346571789).
- XVIII. Derevianko, Ivan. "BLAZOR - IT'S TIME TO FORGET JAVASCRIPT" *Ivan Derevianko*, Ivan Derevianko, 16 July 2019, <https://ivanderevianko.com/2019/07/blazor-its-time-to-forget-javascript>.
- XIX. Szegedy, Christian, et al. "Rethinking the inception architecture for computer vision." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016, <<https://arxiv.org/pdf/1512.00567>>.
- XX. "Ultralytics." *Ultralytics YOLO Docs*, <https://docs.ultralytics.com/>. Accessed 30 Oct. 2024.
- XXI. Pandey, Abhishek Kumar. "SiLU (Sigmoid Linear Unit) activation function." *Medium*, Medium, 6 Apr. 2023, <https://medium.com/@akp83540/silu-sigmoid-linear-unit-activation-function-d9b6845f0c81>.
- XXII. Hidayatullah, Priyanto. "YOLOv8 Architecture Detailed Explanation- A Complete Breakdown" *Youtube*, Dr. Priyanto Hidayatullah, 28 Oct. 2023, [https://www.youtube.com/watch?v=HQXhDO7COj8&t=74s&ab\\_channel=Dr.PriyantoHidayatullah](https://www.youtube.com/watch?v=HQXhDO7COj8&t=74s&ab_channel=Dr.PriyantoHidayatullah).

# References

- XXIII. Terven, Juan, Diana-Margarita Córdova-Esparza, and Julio-Alejandro Romero-González. "A comprehensive review of yolo architectures in computer vision: From yolov1 to yolov8 and yolo-nas." *Machine Learning and Knowledge Extraction* 5.4 (2023): 1680-1716, accessed 29 October 2024, <<https://arxiv.org/pdf/2304.00501v6>>.
- XXIV. Loshchilov, I. "Decoupled weight decay regularization." *arXiv preprint arXiv:1711.05101* (2017), accessed 29 October 2024, <<https://arxiv.org/pdf/1711.05101>>.
- XXV. Pytorch. "Vision/Torchvision/Models/Alexnet.Py at Main · Pytorch/Vision." *GitHub*, [github.com/pytorch/vision/blob/main/torchvision/models/alexnet.py](https://github.com/pytorch/vision/blob/main/torchvision/models/alexnet.py). Accessed 20 Nov. 2024.
- XXVI. Pytorch. "Vision/Torchvision/Models/ResNet.Py at Main · Pytorch/Vision." *GitHub*, <https://github.com/pytorch/vision/blob/main/torchvision/models/resnet.py>. Accessed 20 Nov. 2024.
- XXVII. Pytorch. "Vision/Torchvision/Models/VGG.Py at Main · Pytorch/Vision." *GitHub*, <https://github.com/pytorch/vision/blob/main/torchvision/models/vgg.py>. Accessed 20 Nov. 2024.
- XXVIII. Nvidia. "DeepLearningExamples/Pytorch/Classification/Convnets/Efficientnet at Master · Nvidia/Deeplearningexamples." *GitHub*, [github.com/NVIDIA/DeepLearningExamples/tree/master/PyTorch/Classification/ConvNets/efficientnet](https://github.com/NVIDIA/DeepLearningExamples/tree/master/PyTorch/Classification/ConvNets/efficientnet). Accessed 20 Nov. 2024.
- XXIX. Pytorch. "Vision/Torchvision/Models/Inception.Py at Main · Pytorch/Vision." *GitHub*, <https://github.com/pytorch/vision/blob/main/torchvision/models/inception.py>. Accessed 20 Nov. 2024.